

# Nokia-vs-Hammer Classification

March 2, 2018

```
In [46]: import keras
         from keras.datasets import mnist
         from keras.models import Sequential
         from keras.layers import Dense, Dropout
         from keras.optimizers import RMSprop
         from skimage import io
         import os
         import numpy as np
```

## 0.0.1 hyperparameters

```
In [91]: batch_size = 5
         num_classes = 2
         epochs = 5
```

## 0.0.2 train, test data

```
In [92]: def read_images(path, label=1):
         data = []
         for img in os.listdir(path):
             data.append(io.imread(path + '/' + img, as_grey=True).reshape(1024))
         return data, [label] * len(data)
```

```
In [93]: def compile_set(path, label):
         x_train, y_train = read_images(path + '/train', label)
         x_dev, y_dev = read_images(path + '/dev', label)
         x_test, y_test = read_images(path + '/test', label)
         return {'train': (x_train, y_train),
                 'dev': (x_dev, y_dev),
                 'test': (x_test, y_test)}
```

```
In [94]: # the data, split between train and test sets
         phone_set = compile_set('phone_dataset', label=0)
         hammer_set = compile_set('hammer_dataset', label=1)

         x_train = np.array(phone_set['train'][0] + hammer_set['train'][0])
         y_train = np.array(phone_set['train'][1] + hammer_set['train'][1])
```

```

x_dev = np.array(phone_set['dev'][0] + hammer_set['dev'][0])
y_dev = np.array(phone_set['dev'][1] + hammer_set['dev'][1])

x_test = np.array(phone_set['test'][0] + hammer_set['test'][0])
y_test = np.array(phone_set['test'][1] + hammer_set['test'][1])

print(x_train.shape[0], 'train samples')
print(x_dev.shape[0], 'development samples')
print(x_test.shape[0], 'test samples')

```

```

20 train samples
4 development samples
4 test samples

```

### 0.0.3 neural network classifier

```

In [96]: model = Sequential()
        model.add(Dense(64, activation='relu', input_shape=(1024,)))
        model.add(Dropout(0.5))
        model.add(Dense(1, activation='sigmoid'))

        model.summary()

        model.compile(loss='binary_crossentropy',
                      optimizer=RMSprop(),
                      metrics=['accuracy'])

```

```

-----
Layer (type)                Output Shape                Param #
-----
dense_22 (Dense)             (None, 64)                  65600
-----
dropout_12 (Dropout)         (None, 64)                  0
-----
dense_23 (Dense)             (None, 1)                   65
=====
Total params: 65,665
Trainable params: 65,665
Non-trainable params: 0
-----

```

### 0.0.4 training

```

In [97]: history = model.fit(x_train, y_train,
                             batch_size=batch_size,
                             epochs=epochs,

```

```
verbose=1,  
validation_data=(x_dev, y_dev))
```

Train on 20 samples, validate on 4 samples

Epoch 1/5

20/20 [=====] - 1s 44ms/step - loss: 7.4974 - acc: 0.5000 - val\_loss: 7.9711

Epoch 2/5

20/20 [=====] - 0s 403us/step - loss: 9.6052 - acc: 0.4000 - val\_loss: 7.9711

Epoch 3/5

20/20 [=====] - 0s 822us/step - loss: 7.9784 - acc: 0.5000 - val\_loss: 7.9711

Epoch 4/5

20/20 [=====] - 0s 515us/step - loss: 7.9721 - acc: 0.5000 - val\_loss: 7.9711

Epoch 5/5

20/20 [=====] - 0s 346us/step - loss: 7.9849 - acc: 0.5000 - val\_loss: 7.9711

```
In [98]: score = model.evaluate(x_dev, y_dev, verbose=0)  
         print('Dev loss:', score[0])  
         print('Dev accuracy:', score[1])
```

Dev loss: 7.971192359924316

Dev accuracy: 0.5

### 0.0.5 accuracy

```
In [99]: score = model.evaluate(x_test, y_test, verbose=0)  
         print('Test loss:', score[0])  
         print('Test accuracy:', score[1])
```

Test loss: 7.971192359924316

Test accuracy: 0.5